

Assignment 9: DonorsChoose_RF_GBDT

Data splitting and pre-processing

In [0]:

```
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

import pickle
from tqdm import tqdm
import os
from collections import Counter
```

In [4]:

```
from google.colab import drive
drive.mount("/content/drive")

project_data = pd.read_csv('/content/drive/My Drive/train_data.csv')
resource_data = pd.read_csv('/content/drive/My Drive/resources.csv')
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3Aietf%3Awg%3Aoauth%3A2.O%b&scope=email%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdocs.test%20https%3A%2F%2Fwww.googleapis.com%2Fdrive%20https%3A%2F%2Fwww.googleapis.com%2Fphotos.readonly%20https%3A%2F%2Fwww.googleapis.com%2Fpeopleapi.readonly&response_type=code

Enter your authorization code:
.....

Mounted at /content/drive

In [5]:

```
project_data.isnull().sum()
```

Out[5]:

Unnamed: 0	0
id	0
teacher_id	0
teacher prefix	3

```

school_state                                0
project_submitted_datetime                  0
project_grade_category                      0
project_subject_categories                  0
project_subject_subcategories               0
project_title                              0
project_essay_1                             0
project_essay_2                             0
project_essay_3                           105490
project_essay_4                           105490
project_resource_summary                    0
teacher_number_of_previously_posted_projects 0
project_is_approved                        0
dtype: int64

```

In [6]:

```

#filling 3 null teacher prefix values with Teacher

project_data["teacher_prefix"].fillna("Teacher",inplace = True)
project_data.isnull().sum()

```

Out[6]:

```

Unnamed: 0                                0
id                                          0
teacher_id                                0
teacher_prefix                            0
school_state                              0
project_submitted_datetime                 0
project_grade_category                     0
project_subject_categories                 0
project_subject_subcategories              0
project_title                             0
project_essay_1                           0
project_essay_2                           0
project_essay_3                           0
project_essay_4                           0
project_essay_3                           105490
project_essay_4                           105490
project_resource_summary                   0
teacher_number_of_previously_posted_projects 0
project_is_approved                       0
dtype: int64

```

In [0]:

```

# merge two column text dataframe:
project_data["essay"] = project_data["project_essay_1"].map(str) + \
    project_data["project_essay_2"].map(str) + \
    project_data["project_essay_3"].map(str) + \
    project_data["project_essay_4"].map(str)

```

In [0]:

```

price_data = resource_data.groupby('id').agg({'price':'sum', 'quantity':'sum'}).reset_index()
project_data = pd.merge(project_data, price_data, on='id', how='left')

```

In [9]:

```

project_data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 109248 entries, 0 to 109247
Data columns (total 20 columns):
Unnamed: 0                                109248 non-null int64
id                                          109248 non-null object
teacher_id                                109248 non-null object
teacher_prefix                            109248 non-null object
school_state                              109248 non-null object
project_submitted_datetime                 109248 non-null object
project_grade_category                     109248 non-null object
project_subject_categories                 109248 non-null object
project_subject_subcategories              109248 non-null object

```

```

project_title      109248 non-null object
project_essay_1    109248 non-null object
project_essay_2    109248 non-null object
project_essay_3    3758 non-null object
project_essay_4    3758 non-null object
project_resource_summary 109248 non-null object
teacher_number_of_previously_posted_projects 109248 non-null int64
project_is_approved 109248 non-null int64
essay              109248 non-null object
price              109248 non-null float64
quantity           109248 non-null int64
dtypes: float64(1), int64(4), object(15)
memory usage: 17.5+ MB

```

Considering only 20k train , 5k validation and 5k test data points

In [0]:

```

from sklearn.utils import resample
p_d = resample(project_data,n_samples = 30000)

```

In [0]:

```

#splitting data as 30% to test
y = p_d["project_is_approved"]
X = p_d
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.16, random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.19, random_state=42)
)

```

In [12]:

```

print(X_train.shape," ",y_train.shape)
print(X_test.shape," ",y_test.shape)
print(X_val.shape," ",y_val.shape)

```

```

(20412, 20)    (20412,)
(4800, 20)     (4800,)
(4788, 20)     (4788,)

```

Function for Response coding

In [0]:

```

# function for response coding

def get_gv_fea_dict(alpha, feature, df):
    value_count = X_train[feature].value_counts()
    gv_dict = dict()

    for i, denominator in value_count.items():
        vec = []
        for k in range(2):
            cls_cnt = X_train.loc[(X_train['project_is_approved']==k) & (X_train[feature]==i)]
            vec.append((cls_cnt.shape[0] + alpha*10)/ (denominator + 20*alpha))

        # we are adding the gene/variation to the dict as key and vec as value
        gv_dict[i]=vec
    return gv_dict

def response_encode(alpha, feature, df):
    """
    This function calculates probability of category belonging to class 1.
    As there are only two classes. so  $P(y=0/C) = 1 - P(y=1/c)$ 

    Returns dictionary of # elements in positive class and their probabilities.

    """

```

```

gv_dict = get_gv_fea_dict(alpha, feature, df)
value_count = X_train[feature].value_counts()

gv_fea = []
for index, row in df.iterrows():
    if row[feature] in dict(value_count).keys():
        gv_fea.append(gv_dict[row[feature]])
    else:
        gv_fea.append([1/2, 1/2])
return gv_fea

```

Preprocessing categorical Features

1. project subject categories

In [0]:

```

#using code from assignment
# project subject categories
categories = list(X_train['project_subject_categories'].values)

cat_list = []
for i in categories:
    temp = ""
    for j in i.split(','):
        if 'The' in j.split():
            j=j.replace('The','')
        j = j.replace(' ','')
        temp+=j.strip()+" "
        temp = temp.replace('&','_')
    cat_list.append(temp.strip())

X_train['clean_categories'] = cat_list
X_train.drop(['project_subject_categories'], axis=1, inplace=True)

from collections import Counter
my_counter = Counter()
for word in X_train['clean_categories'].values:
    my_counter.update(word.split())

cat_dict = dict(my_counter)
sorted_cat_dict = dict(sorted(cat_dict.items(), key=lambda kv: kv[1]))

# project subject categories for test data

categories = list(X_test['project_subject_categories'].values)

cat_list = []
for i in categories:
    temp = ""
    for j in i.split(','):
        if 'The' in j.split():
            j=j.replace('The','')
        j = j.replace(' ','')
        temp+=j.strip()+" "
        temp = temp.replace('&','_')
    cat_list.append(temp.strip())

X_test['clean_categories'] = cat_list
X_test.drop(['project_subject_categories'], axis=1, inplace=True)

# project subject categories for test data

categories = list(X_val['project_subject_categories'].values)

cat_list = []
for i in categories:
    temp = ""
    for j in i.split(','):
        if 'The' in j.split():
            j=j.replace('The','')
        j = j.replace(' ','')

```

```

        temp+=j.strip()+" "
        temp = temp.replace('&','_')
        cat_list.append(temp.strip())

X_val['clean_categories'] = cat_list
X_val.drop(['project_subject_categories'], axis=1, inplace=True)

```

In [0]:

```

alpha = 1
categories_feature_responseCoding = np.array(response_encode(alpha, "clean_categories", X_train))
val_categories_feature_responseCoding = np.array(response_encode(alpha, "clean_categories",
X_val))
test_categories_feature_responseCoding = np.array(response_encode(alpha, "clean_categories",
X_test))

```

1. project subject sub_categories

In [0]:

```

sub_categories = list(X_train['project_subject_subcategories'].values)
sub_cat_list = []
for i in sub_categories:
    temp = ""
    for j in i.split(','):
        if 'The' in j.split(): # this will split each of the category based on space "Math & Science"
            e=> "Math","&", "Science"
            j=j.replace('The','') # if we have the words "The" we are going to replace it with ''(i.e removing 'The')
            j = j.replace(' ','')
            temp +=j.strip()+" "
            temp = temp.replace('&','_')
        sub_cat_list.append(temp.strip())

X_train['clean_subcategories'] = sub_cat_list
X_train.drop(['project_subject_subcategories'], axis=1, inplace=True)

# count of all the words in corpus python: https://stackoverflow.com/a/22898595/4084039
my_counter = Counter()
for word in X_train['clean_subcategories'].values:
    my_counter.update(word.split())

sub_cat_dict = dict(my_counter)
sorted_sub_cat_dict = dict(sorted(sub_cat_dict.items(), key=lambda kv: kv[1]))

sub_categories = list(X_test['project_subject_subcategories'].values)
sub_cat_list = []
for i in sub_categories:
    temp = ""
    for j in i.split(','):
        if 'The' in j.split(): # this will split each of the category based on space "Math & Science"
            e=> "Math","&", "Science"
            j=j.replace('The','') # if we have the words "The" we are going to replace it with ''(i.e removing 'The')
            j = j.replace(' ','')
            temp +=j.strip()+" "
            temp = temp.replace('&','_')
        sub_cat_list.append(temp.strip())

X_test['clean_subcategories'] = sub_cat_list
X_test.drop(['project_subject_subcategories'], axis=1, inplace=True)

# for validation data set
sub_categories = list(X_val['project_subject_subcategories'].values)
sub_cat_list = []
for i in sub_categories:
    temp = ""
    for j in i.split(','):
        if 'The' in j.split(): # this will split each of the category based on space "Math & Science"
            e=> "Math","&", "Science"
            j=j.replace('The','') # if we have the words "The" we are going to replace it with ''(i.e removing 'The')

```

```

        j = j.replace(' ', '')
        temp += j.strip() + " "
        temp = temp.replace('&', '_')
        sub_cat_list.append(temp.strip())

```

```

X_val['clean_subcategories'] = sub_cat_list
X_val.drop(['project_subject_subcategories'], axis=1, inplace=True)

```

In [0]:

```

alpha = 1
subcategories_feature_responseCoding = np.array(response_encode(alpha, "clean_subcategories",
X_train))
val_subcategories_feature_responseCoding = np.array(response_encode(alpha, "clean_subcategories",
X_val))
test_subcategories_feature_responseCoding = np.array(response_encode(alpha, "clean_subcategories",
X_test))

```

In [18]:

```

len(subcategories_feature_responseCoding)

```

Out[18]:

20412

1. Teacher Prefix

In [19]:

```

#preprocessing teacher prefix
prefix = list(X_train['teacher_prefix'].values)
prefix_list = []
for i in prefix:
    temp = ""
    if "." in i:
        i=i.replace('.', '')
        temp+=i.strip()+" "
    prefix_list.append(temp.strip())

X_train['clean_prefix'] = prefix_list

my_counter = Counter()
for word in X_train['clean_prefix'].values:
    my_counter.update(word.split())

prefix_dict = dict(my_counter)
sorted_prefix_dict = dict(sorted(prefix_dict.items(), key=lambda kv: kv[1]))
print(sorted_prefix_dict)

#preprocessing teacher prefix for test data
prefix = list(X_test['teacher_prefix'].values)
prefix_list = []
for i in prefix:
    temp = ""
    if "." in i:
        i=i.replace('.', '')
        temp+=i.strip()+" "
    prefix_list.append(temp.strip())

X_test['clean_prefix'] = prefix_list

#preprocessing teacher prefix for val data
prefix = list(X_val['teacher_prefix'].values)
prefix_list = []
for i in prefix:
    temp = ""
    if "." in i:
        i=i.replace('.', '')
        temp+=i.strip()+" "
    prefix_list.append(temp.strip())

```

```
X_val['clean_prefix'] = prefix_list
```

```
{'Dr': 2, 'Teacher': 451, 'Mr': 2002, 'Ms': 7262, 'Mrs': 10695}
```

```
In [0]:
```

```
alpha = 1
# train gene feature
prefix_feature_responseCoding = np.array(response_encode(alpha, "clean_prefix", X_train))
val_prefix_feature_responseCoding = np.array(response_encode(alpha, "clean_prefix", X_val))
test_prefix_feature_responseCoding = np.array(response_encode(alpha, "clean_prefix", X_test))
```

1. Project Grade Category

```
In [21]:
```

```
# preprocessing of grade category for train data

grade = list(X_train['project_grade_category'].values)
grade_list = []
for i in grade:
    temp = ""
    if "Grades" in i:
        i = i.replace("Grades", "")
    if "6-8" in i:
        i = i.replace("6-8", "six_eight")
    if "3-5" in i:
        i = i.replace("3-5", "three_five")
    if "9-12" in i:
        i = i.replace("9-12", "nine_twelve")
    if "PreK-2" in i:
        i = i.replace("PreK-2", "prek_two")
    temp+=i.strip()+" "
    grade_list.append(temp.strip())

X_train['clean_grade'] = grade_list

my_counter = Counter()
for word in X_train['clean_grade'].values:
    my_counter.update(word.split())

grade_dict = dict(my_counter)
sorted_grade_dict = dict(sorted(grade_dict.items(), key=lambda kv: kv[1]))
print(sorted_grade_dict)

# preprocessing of grade category for test data

grade = list(X_test['project_grade_category'].values)
grade_list = []
for i in grade:
    temp = ""
    if "Grades" in i:
        i = i.replace("Grades", "")
    if "6-8" in i:
        i = i.replace("6-8", "six_eight")
    if "3-5" in i:
        i = i.replace("3-5", "three_five")
    if "9-12" in i:
        i = i.replace("9-12", "nine_twelve")
    if "PreK-2" in i:
        i = i.replace("PreK-2", "prek_two")
    temp+=i.strip()+" "
    grade_list.append(temp.strip())

X_test['clean_grade'] = grade_list

grade = list(X_val['project_grade_category'].values)
grade_list = []
for i in grade:
    temp = ""
    if "Grades" in i:
```

```

i = i.replace("Grades", "")
if "6-8" in i:
    i = i.replace("6-8", "six_eight")
if "3-5" in i:
    i = i.replace("3-5", "three_five")
if "9-12" in i:
    i = i.replace("9-12", "nine_twelve")
if "PreK-2" in i:
    i = i.replace("PreK-2", "prek_two")
temp+=i.strip()+" "
grade_list.append(temp.strip())

```

```
X_val['clean_grade'] = grade_list
```

```
{'nine_twelve': 2056, 'six_eight': 3192, 'three_five': 7027, 'prek_two': 8137}
```

In [0]:

```

alpha = 1
# train gene feature
grade_feature_responseCoding = np.array(response_encode(alpha, "clean_grade", X_train))
val_grade_feature_responseCoding = np.array(response_encode(alpha, "clean_grade", X_val))
test_grade_feature_responseCoding = np.array(response_encode(alpha, "clean_grade", X_test))

```

1. School State

In [0]:

```

#no need of preprocessing on school state

state = X_train["school_state"].value_counts()
sorted_state = dict(state)
sorted_state_dict = dict(sorted(sorted_state.items(), key=lambda kv: kv[1]))
X_train["clean_state"] = X_train["school_state"]

#similarly for X_test
X_test["clean_state"] = X_test["school_state"]

#similarly for X_val
X_val["clean_state"] = X_val["school_state"]

```

In [0]:

```

alpha = 1
# train gene feature
state_feature_responseCoding = np.array(response_encode(alpha, "clean_state", X_train))
val_state_feature_responseCoding = np.array(response_encode(alpha, "clean_state", X_val))
test_state_feature_responseCoding = np.array(response_encode(alpha, "clean_state", X_test))

```

Preprocessing Numerical Feature

1. Standardizing price

In [25]:

```

from sklearn.preprocessing import StandardScaler

price_scalar = StandardScaler()
price_scalar.fit(project_data['price'].values.reshape(-1,1))
print(f"Mean : {price_scalar.mean_[0]}, Standard deviation : {np.sqrt(price_scalar.var_[0])}")

#train data price standardization
price_standardized = price_scalar.transform(X_train['price'].values.reshape(-1, 1))

#val data price stanardization. Fit method applied on X_train
val_price_standardized = price_scalar.transform(X_val['price'].values.reshape(-1, 1))

```



```
#test data price stanardization. Fit method applied on X_train
test_price_standardized = price_scalar.transform(X_test['price'].values.reshape(-1, 1))
```

Mean : 298.1193425966608, Standard deviation : 367.49634838483496

1. Standardizing quantity

In [26]:

```
price_scalar = StandardScaler()
price_scalar.fit(X_train["quantity"].values.reshape(-1, 1))
print(f"Mean of Quantity : {price_scalar.mean_[0]}, Standard deviation of Quantity :
{np.sqrt(price_scalar.var_[0])}")

#train data quantity stanardization
quantity_standardized = price_scalar.transform(X_train["quantity"].values.reshape(-1, 1))

#val data quantity stanardization. Fit method applied on X_train
val_quantity_standardized = price_scalar.transform(X_val["quantity"].values.reshape(-1, 1))

#test data quantity stanardization. Fit method applied on X_train
test_quantity_standardized = price_scalar.transform(X_test["quantity"].values.reshape(-1, 1))
```

Mean of Quantity : 17.086909660983736, Standard deviation of Quantity : 26.07678878802128

1. Standardizing number of ppp

In [27]:

```
price_scalar = StandardScaler()
price_scalar.fit(X_train['teacher_number_of_previously_posted_projects'].values.reshape(-1,1))
print(f"Mean : {price_scalar.mean_[0]}, Standard deviation : {np.sqrt(price_scalar.var_[0])}")

#train data ppp standardization
number_ppp_standardized =
price_scalar.transform(X_train['teacher_number_of_previously_posted_projects'].values.reshape(-1,
1))

#val data price stanardization. Fit method applied on X_train
val_number_ppp_standardized =
price_scalar.transform(X_val['teacher_number_of_previously_posted_projects'].values.reshape(-1,
1))

#test data price stanardization. Fit method applied on X_train
test_number_ppp_standardized =
price_scalar.transform(X_test['teacher_number_of_previously_posted_projects'].values.reshape(-1, 1)
)
```

Mean : 10.909122085048011, Standard deviation : 26.76311411192372

Preprocessing of Text Feature for both test and train data

In [0]:

```
#using function and stopwords form assignemnt

import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
```

```

phrase = re.sub(r"\\'ll", " will", phrase)
phrase = re.sub(r"\\'t", " not", phrase)
phrase = re.sub(r"\\'ve", " have", phrase)
phrase = re.sub(r"\\'m", " am", phrase)
return phrase

# we are removing the words from the stop words list: 'no', 'nor', 'not'
stopwords= ['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've",
\
            "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his',
'himself', \
            'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'them',
'their', \
            'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll",
'these', 'those', \
            'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having',
'do', 'does', \
            'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as', 'until', '
while', 'of', \
            'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through', 'during',
'before', 'after', \
            'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over', 'under'
, 'again', 'further', \
            'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any', 'both', 'e
ach', 'few', 'more', \
            'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too', 'very', \
            's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'now', 'd', 'll'
, 'm', 'o', 're', \
            've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn', "dc
esn't", 'hadn', \
            "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn',
"mightn't", 'mustn', \
            'mustn't', 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'wasn',
"wasn't", 'weren', "weren't", \
            'won', "won't", 'wouldn', "wouldn't"]

```

1. preprocessing of project essay

In [29]:

```

from tqdm import tqdm

#for train data
preprocessed_essays = []
# tqdm is for printing the status bar
for sentence in tqdm(X_train['essay'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\\r', ' ')
    sent = sent.replace('\\n', ' ')
    sent = sent.replace('\\n', ' ')
    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    # https://gist.github.com/sebleier/554280
    sent = ' '.join(e for e in sent.split() if e not in stopwords)
    preprocessed_essays.append(sent.lower().strip())

val_preprocessed_essays = []
# tqdm is for printing the status bar
for sentence in tqdm(X_val['essay'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\\r', ' ')
    sent = sent.replace('\\n', ' ')
    sent = sent.replace('\\n', ' ')
    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    # https://gist.github.com/sebleier/554280
    sent = ' '.join(e for e in sent.split() if e not in stopwords)
    val_preprocessed_essays.append(sent.lower().strip())

test_preprocessed_essays = []
# tqdm is for printing the status bar
for sentence in tqdm(X_test['essay'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\\r', ' ')
    sent = sent.replace('\\n', ' ')

```

```

sent = sent.replace('\n', ' ')
sent = sent.replace('\n', ' ')
sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
# https://gist.github.com/sebleier/554280
sent = ' '.join(e for e in sent.split() if e not in stopwords)
test_preprocessed_essays.append(sent.lower().strip())

```

```

100%|██████████| 20412/20412 [00:12<00:00, 1620.78it/s]
100%|██████████| 4788/4788 [00:03<00:00, 1546.50it/s]
100%|██████████| 4800/4800 [00:03<00:00, 1573.16it/s]

```

1. preprocessing of project title

In [30]:

```

preprocessed_title = []
# tqdm is for printing the status bar
for sentence in tqdm(X_train['project_title'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\r', ' ')
    sent = sent.replace('\n', ' ')
    sent = sent.replace('\n', ' ')
    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    # https://gist.github.com/sebleier/554280
    sent = ' '.join(e for e in sent.split() if e not in stopwords)
    preprocessed_title.append(sent.lower().strip())

# for val data
val_preprocessed_title = []
# tqdm is for printing the status bar
for sentence in tqdm(X_val['project_title'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\r', ' ')
    sent = sent.replace('\n', ' ')
    sent = sent.replace('\n', ' ')
    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    # https://gist.github.com/sebleier/554280
    sent = ' '.join(e for e in sent.split() if e not in stopwords)
    val_preprocessed_title.append(sent.lower().strip())

# for test data
test_preprocessed_title = []
# tqdm is for printing the status bar
for sentence in tqdm(X_test['project_title'].values):
    sent = decontracted(sentence)
    sent = sent.replace('\r', ' ')
    sent = sent.replace('\n', ' ')
    sent = sent.replace('\n', ' ')
    sent = re.sub('[^A-Za-z0-9]+', ' ', sent)
    # https://gist.github.com/sebleier/554280
    sent = ' '.join(e for e in sent.split() if e not in stopwords)
    test_preprocessed_title.append(sent.lower().strip())

```

```

100%|██████████| 20412/20412 [00:00<00:00, 33519.32it/s]
100%|██████████| 4788/4788 [00:00<00:00, 32927.41it/s]
100%|██████████| 4800/4800 [00:00<00:00, 33268.16it/s]

```

Vectorizing Text Feature

1. BOW

In [31]:

```

vectorizer = CountVectorizer(min_df=10,ngram_range=(1,4),max_features=5000)
#fit using train data
vectorizer.fit(preprocessed_essays)
essay_feature = vectorizer.get_feature_names()

# for train data

```

```

# for train data
text_bow = vectorizer.transform(preprocessed_essays)
print("Shape of train matrix : ",text_bow.shape)

# for val data
val_text_bow = vectorizer.transform(val_preprocessed_essays)

# for test data
test_text_bow = vectorizer.transform(test_preprocessed_essays)
print("Shape of test matrix : ",test_text_bow.shape)

# for title
vectorizer.fit(preprocessed_title)
title_feature = vectorizer.get_feature_names()

# for train data
title_bow = vectorizer.transform(preprocessed_title)
print("Shape of train matrix : ",title_bow.shape)

# for val data
val_title_bow = vectorizer.transform(val_preprocessed_title)
print("Shape of test matrix : ",val_title_bow.shape)

# for test data
test_title_bow = vectorizer.transform(test_preprocessed_title)
print("Shape of test matrix : ",test_title_bow.shape)

```

```

Shape of train matrix : (20412, 5000)
Shape of test matrix : (4800, 5000)
Shape of train matrix : (20412, 1824)
Shape of test matrix : (4788, 1824)
Shape of test matrix : (4800, 1824)

```

1. TFIDF

In [32]:

```

vectorizer = TfidfVectorizer(min_df=10,ngram_range=(1,4),max_features=5000)
#fit using train data
vectorizer.fit(preprocessed_essays)
essay_feature_tfidf = vectorizer.get_feature_names()
# for train data
text_tfidf = vectorizer.transform(preprocessed_essays)
print("Shape of train matrix : ",text_tfidf.shape)

# for val data
val_text_tfidf = vectorizer.transform(val_preprocessed_essays)
print("Shape of val matrix : ",val_text_tfidf.shape)

# for test data
test_text_tfidf = vectorizer.transform(test_preprocessed_essays)
print("Shape of test matrix : ",test_text_tfidf.shape)

# for title
vectorizer.fit(preprocessed_title)
title_feature_tfidf = vectorizer.get_feature_names()

# for train data
title_tfidf = vectorizer.transform(preprocessed_title)
print("Shape of train matrix : ",title_tfidf.shape)

# for val data
val_title_tfidf = vectorizer.transform(val_preprocessed_title)
print("Shape of val matrix : ",val_title_tfidf.shape)

# for test data
test_title_tfidf = vectorizer.transform(test_preprocessed_title)
print("Shape of test matrix : ",test_title_tfidf.shape)

```

```
Shape of train matrix : (20412, 5000)
Shape of val matrix : (4788, 5000)
Shape of test matrix : (4800, 5000)
Shape of train matrix : (20412, 1824)
Shape of val matrix : (4788, 1824)
Shape of test matrix : (4800, 1824)
```

1. Avg W2v

In [0]:

```
with open('/content/drive/My Drive/glove_vectors', 'rb') as f:
    model = pickle.load(f)
    glove_words = set(model.keys())
```

In [34]:

```
# for train data
avg_w2v_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sentence in tqdm(preprocessed_essays): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    avg_w2v_vectors.append(vector)
print(len(avg_w2v_vectors))
print(len(avg_w2v_vectors[0]))

# for val data
val_avg_w2v_vectors = [] # the avg-w2v for each sentence/review is stored in this list
for sentence in tqdm(val_preprocessed_essays): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    val_avg_w2v_vectors.append(vector)

print(len(val_avg_w2v_vectors))
print(len(val_avg_w2v_vectors[0]))

# for test data
test_avg_w2v_vectors = [] # the avg-w2v for each sentence/review is stored in this list
for sentence in tqdm(test_preprocessed_essays): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    test_avg_w2v_vectors.append(vector)

print(len(test_avg_w2v_vectors))
print(len(test_avg_w2v_vectors[0]))
```

```
100%|██████████| 20412/20412 [00:07<00:00, 2790.82it/s]
 6%|███████| 279/4788 [00:00<00:01, 2780.85it/s]
```

20412
300

100%|██████████| 4788/4788 [00:01<00:00, 2815.25it/s]
12%|███| 552/4800 [00:00<00:01, 2772.16it/s]

4788
300

100%|██████████| 4800/4800 [00:01<00:00, 2778.64it/s]

4800
300

In [35]:

```
title_avg_w2v_vectors = []
for sentence in tqdm(preprocessed_title):
    vector = np.zeros(300)
    cnt_words = 0;
    for word in sentence.split():
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    title_avg_w2v_vectors.append(vector)

print(len(title_avg_w2v_vectors))
print(len(title_avg_w2v_vectors[0]))

# for val data
val_title_avg_w2v_vectors = []
for sentence in tqdm(val_preprocessed_title):
    vector = np.zeros(300)
    cnt_words = 0;
    for word in sentence.split():
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    val_title_avg_w2v_vectors.append(vector)

print(len(val_title_avg_w2v_vectors))
print(len(val_title_avg_w2v_vectors[0]))

# for test data
test_title_avg_w2v_vectors = []
for sentence in tqdm(test_preprocessed_title):
    vector = np.zeros(300)
    cnt_words = 0;
    for word in sentence.split():
        if word in glove_words:
            vector += model[word]
            cnt_words += 1
    if cnt_words != 0:
        vector /= cnt_words
    test_title_avg_w2v_vectors.append(vector)

print(len(test_title_avg_w2v_vectors))
print(len(test_title_avg_w2v_vectors[0]))
```

100%|██████████| 20412/20412 [00:00<00:00, 61666.90it/s]
100%|██████████| 4788/4788 [00:00<00:00, 57053.69it/s]
100%|██████████| 4800/4800 [00:00<00:00, 58946.36it/s]

20412
300
4788
300
4800
300

1. TFIDF avgw2v

In [36]:

```
# for train data
tfidf_model = TfidfVectorizer()
tfidf_model.fit(preprocessed_essays)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(tfidf_model.get_feature_names(), list(tfidf_model.idf_)))
tfidf_words = set(tfidf_model.get_feature_names())

tfidf_w2v_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sentence in tqdm(preprocessed_essays): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in tfidf_words):
            vec = model[word] # getting the vector for each word
            # here we are multiplying idf value(dictionary[word]) and the tf
            value((sentence.count(word)/len(sentence.split())))
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) # getting the tf
            idf value for each word
            vector += (vec * tf_idf) # calculating tfidf weighted w2v
            tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    tfidf_w2v_vectors.append(vector)

print(len(tfidf_w2v_vectors))
print(len(tfidf_w2v_vectors[0]))

# ----- for val data -----
val_tfidf_model = TfidfVectorizer()
val_tfidf_model.fit(preprocessed_essays)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(val_tfidf_model.get_feature_names(), list(val_tfidf_model.idf_)))
tfidf_words = set(val_tfidf_model.get_feature_names())

val_tfidf_w2v_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sentence in tqdm(val_preprocessed_essays): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in tfidf_words):
            vec = model[word] # getting the vector for each word
            # here we are multiplying idf value(dictionary[word]) and the tf
            value((sentence.count(word)/len(sentence.split())))
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) # getting the tf
            idf value for each word
            vector += (vec * tf_idf) # calculating tfidf weighted w2v
            tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    val_tfidf_w2v_vectors.append(vector)

print(len(val_tfidf_w2v_vectors))
print(len(val_tfidf_w2v_vectors[0]))

# ----- for test data -----
test_tfidf_model = TfidfVectorizer()
test_tfidf_model.fit(preprocessed_essays)
# we are converting a dictionary with word as a key, and the idf as a value
dictionary = dict(zip(test_tfidf_model.get_feature_names(), list(test_tfidf_model.idf_)))
tfidf_words = set(test_tfidf_model.get_feature_names())

test_tfidf_w2v_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sentence in tqdm(test_preprocessed_essays): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf idf weight = 0; # num of words with a valid vector in the sentence/review
```

```

for word in sentence.split(): # for each word in a review/sentence
    if (word in glove_words) and (word in tfidf_words):
        vec = model[word] # getting the vector for each word
        # here we are multiplying idf value(dictionary[word]) and the tf
value((sentence.count(word)/len(sentence.split())))
        tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) # getting the tf
idf value for each word
        vector += (vec * tf_idf) # calculating tfidf weighted w2v
        tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    test_tfidf_w2v_vectors.append(vector)

print(len(test_tfidf_w2v_vectors))
print(len(test_tfidf_w2v_vectors[0]))

```

100%|██████████| 20412/20412 [00:42<00:00, 476.18it/s]

20412
300

100%|██████████| 4788/4788 [00:10<00:00, 472.08it/s]

4788
300

100%|██████████| 4800/4800 [00:10<00:00, 465.20it/s]

4800
300

In [37]:

```

# for title
tfidf_model.fit(preprocessed_title)

dictionary = dict(zip(tfidf_model.get_feature_names(), list(tfidf_model.idf_)))
tfidf_words = set(tfidf_model.get_feature_names())

title_tfidf_w2v_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sentence in tqdm(preprocessed_title): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in tfidf_words):
            vec = model[word] # getting the vector for each word
            # here we are multiplying idf value(dictionary[word]) and the tf
value((sentence.count(word)/len(sentence.split())))
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) # getting the tf
idf value for each word
            vector += (vec * tf_idf) # calculating tfidf weighted w2v
            tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    title_tfidf_w2v_vectors.append(vector)

print(len(title_tfidf_w2v_vectors))

# ----- for val title -----
val_tfidf_model.fit(preprocessed_title)

dictionary = dict(zip(val_tfidf_model.get_feature_names(), list(val_tfidf_model.idf_)))
tfidf_words = set(val_tfidf_model.get_feature_names())

val_title_tfidf_w2v_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sentence in tqdm(val_preprocessed_title): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review

```



```

for word in sentence.split(): # for each word in a review/sentence
    if (word in glove_words) and (word in tfidf_words):
        vec = model[word] # getting the vector for each word
        # here we are multiplying idf value(dictionary[word]) and the tf
        value((sentence.count(word)/len(sentence.split())))
        tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) # getting the tf
        idf value for each word
        vector += (vec * tf_idf) # calculating tfidf weighted w2v
        tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    val_title_tfidf_w2v_vectors.append(vector)

print(len(val_title_tfidf_w2v_vectors))

# ----- for test title -----
test_tfidf_model.fit(preprocessed_title)

dictionary = dict(zip(test_tfidf_model.get_feature_names(), list(test_tfidf_model.idf_)))
tfidf_words = set(test_tfidf_model.get_feature_names())

test_title_tfidf_w2v_vectors = []; # the avg-w2v for each sentence/review is stored in this list
for sentence in tqdm(test_preprocessed_title): # for each review/sentence
    vector = np.zeros(300) # as word vectors are of zero length
    tf_idf_weight = 0; # num of words with a valid vector in the sentence/review
    for word in sentence.split(): # for each word in a review/sentence
        if (word in glove_words) and (word in tfidf_words):
            vec = model[word] # getting the vector for each word
            # here we are multiplying idf value(dictionary[word]) and the tf
            value((sentence.count(word)/len(sentence.split())))
            tf_idf = dictionary[word]*(sentence.count(word)/len(sentence.split())) # getting the tf
            idf value for each word
            vector += (vec * tf_idf) # calculating tfidf weighted w2v
            tf_idf_weight += tf_idf
    if tf_idf_weight != 0:
        vector /= tf_idf_weight
    test_title_tfidf_w2v_vectors.append(vector)

print(len(test_title_tfidf_w2v_vectors))

```

```

100%|██████████| 20412/20412 [00:00<00:00, 22591.59it/s]
0%|          | 0/4788 [00:00<?, ?it/s]

```

20412

```

100%|██████████| 4788/4788 [00:00<00:00, 19926.17it/s]
0%|          | 0/4800 [00:00<?, ?it/s]

```

4788

```

100%|██████████| 4800/4800 [00:00<00:00, 18741.97it/s]

```

4800

Printing all

In [38]:

```

print("Text Features that are considered :- ")
print("*"*70)
print("Project Essay BOW:- ",text_bow.shape)
print("Project Essay TFIDF:- ",text_tfidf.shape)
print("*"*70)
print("Project Title BOW:- ",title_bow.shape)
print("Project Title TFIDF:- ",title_tfidf.shape)
print("*"*70)

```

```

Text Features that are considered :-
*****
Project Essay BOW:- (20412, 5000)
Project Essay TFIDF:- (20412, 5000)
*****
Project Title BOW:- (20412, 1824)
Project Title TFIDF:- (20412, 1824)
*****

```

In [0]:

```

avg_w2v_vectors = np.array(avg_w2v_vectors)
title_avg_w2v_vectors = np.array(title_avg_w2v_vectors)
val_avg_w2v_vectors = np.array(val_avg_w2v_vectors)

```

In [0]:

```

test_avg_w2v_vectors = np.array(test_avg_w2v_vectors)
test_title_avg_w2v_vectors = np.array(test_title_avg_w2v_vectors)

```

sets

In [41]:

```

#combining all feature into one
from scipy.sparse import hstack

set1 =
hstack((categories_feature_responseCoding,subcategories_feature_responseCoding,prefix_feature_respo
nseCoding,grade_feature_responseCoding,state_feature_responseCoding,price_standardized,quantity_sta
ndardized,number_ppp_standardized,text_bow,title_bow))
set1_v = hstack((val_categories_feature_responseCoding,val_subcategories_feature_responseCoding,va
l_prefix_feature_responseCoding,val_grade_feature_responseCoding,val_state_feature_responseCoding,
val_price_standardized,val_quantity_standardized,val_number_ppp_standardized,val_text_bow,val_titl
e_bow))
set1_t = hstack((test_categories_feature_responseCoding,test_subcategories_feature_responseCoding,
test_prefix_feature_responseCoding,test_grade_feature_responseCoding,test_state_feature_responseCoc
ng,test_price_standardized,test_quantity_standardized,test_number_ppp_standardized,test_text_bow,t
est_title_bow))

set2 =
hstack((categories_feature_responseCoding,subcategories_feature_responseCoding,prefix_feature_respo
nseCoding,grade_feature_responseCoding,state_feature_responseCoding,text_tfidf,title_tfidf,price_st
andardized,quantity_standardized,number_ppp_standardized))
set2_v = hstack((val_categories_feature_responseCoding,val_subcategories_feature_responseCoding,va
l_prefix_feature_responseCoding,val_grade_feature_responseCoding,val_state_feature_responseCoding,
val_text_tfidf,val_title_tfidf,val_price_standardized,val_quantity_standardized,val_number_ppp_star
dardized))
set2_t = hstack((test_categories_feature_responseCoding,test_subcategories_feature_responseCoding,
test_prefix_feature_responseCoding,test_grade_feature_responseCoding,test_state_feature_responseCoc
ng,test_text_tfidf,test_title_tfidf,test_price_standardized,test_quantity_standardized,test_number_
ppp_standardized))

set3 =
np.hstack((categories_feature_responseCoding,subcategories_feature_responseCoding,prefix_feature_re
sponseCoding,grade_feature_responseCoding,state_feature_responseCoding,price_standardized,quantity_
standardized,number_ppp_standardized,avg_w2v_vectors,title_avg_w2v_vectors))
set3_v = np.hstack((val_categories_feature_responseCoding,val_subcategories_feature_responseCoding
,val_prefix_feature_responseCoding,val_grade_feature_responseCoding,val_state_feature_responseCodir
,val_price_standardized,val_quantity_standardized,val_number_ppp_standardized,val_avg_w2v_vectors,
val_title_avg_w2v_vectors))
set3_t =
np.hstack((test_categories_feature_responseCoding,test_subcategories_feature_responseCoding,test_p
refix_feature_responseCoding,test_grade_feature_responseCoding,test_state_feature_responseCoding,t
est_price_standardized,test_quantity_standardized,test_number_ppp_standardized,test_avg_w2v_vectors
,test_title_avg_w2v_vectors))

set4 =
np.hstack((categories_feature_responseCoding,subcategories_feature_responseCoding,prefix_feature_re
sponseCoding,grade_feature_responseCoding,state_feature_responseCoding,price_standardized,quantity_
standardized,number_ppp_standardized,tfidf_w2v_vectors,title_tfidf_w2v_vectors))
set4_v = np.hstack((val_categories_feature_responseCoding,val_subcategories_feature_responseCoding,
val_prefix_feature_responseCoding,val_grade_feature_responseCoding,val_state_feature_responseCoding,
val_price_standardized,val_quantity_standardized,val_number_ppp_standardized,val_tfidf_w2v_vectors,
val_title_tfidf_w2v_vectors))
set4_t = np.hstack((test_categories_feature_responseCoding,test_subcategories_feature_responseCoding,
test_prefix_feature_responseCoding,test_grade_feature_responseCoding,test_state_feature_responseCoding,
test_price_standardized,test_quantity_standardized,test_number_ppp_standardized,test_tfidf_w2v_vectors,
test_title_tfidf_w2v_vectors))

```

```

set4_v = np.hstack((val_categories_reature_responseCoding, val_subcategories_reature_responseCoding,
, val_prefix_feature_responseCoding, val_grade_feature_responseCoding, val_state_feature_responseCoding,
, val_price_standardized, val_quantity_standardized, val_number_ppp_standardized, val_tfidf_w2v_vectors,
, val_title_tfidf_w2v_vectors))
set4_t =
np.hstack((test_categories_feature_responseCoding, test_subcategories_feature_responseCoding, test_p
refix_feature_responseCoding, test_grade_feature_responseCoding, test_state_feature_responseCoding, t
est_price_standardized, test_quantity_standardized, test_number_ppp_standardized, test_tfidf_w2v_vecto
rs, test_title_tfidf_w2v_vectors))

print(set1.shape, "\t", set1_t.shape)
print(set2.shape, "\t", set2_t.shape)
print(set3.shape, "\t", set3_t.shape)
print(set4.shape, "\t", set4_t.shape)

```

```

(20412, 6837)    (4800, 6837)
(20412, 6837)    (4800, 6837)
(20412, 613)     (4800, 613)
(20412, 613)     (4800, 613)

```

In [0]:

```

from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
from xgboost import XGBClassifier

```

In [0]:

```

def cf_matrix(cm, msg):
    """
    This function is to plot confusion matrix.

    """
    sns.heatmap(cm, annot=True, fmt="d")
    plt.ylabel("Actual Class")
    plt.xlabel("Predicted Class")
    plt.title("Confusion Matrix for {}".format(msg))

```

In [0]:

```

n_estimators = [5, 10, 50, 100, 200, 500, 1000]
max_depth = [2, 3, 4, 5, 6, 7, 8, 9, 10]

def Plot_heatmap(mean_score, name):
    """
    This function plots heatmap.

    """
    df = pd.DataFrame(mean_score, index = n_estimators, columns = max_depth)
    sns.heatmap(df, annot = True)
    plt.ylabel("N_estimators")
    plt.xlabel("Max_depth")
    plt.title(name)
    plt.show()

```

SET1 (BOW)

In [0]:

```

train_auc = []
cv_auc = []
heatmap = []
n_estimators = [5, 10, 50, 100, 200, 500, 1000]
max_depth = [2, 3, 4, 5, 6, 7, 8, 9, 10]

for i in n_estimators:
    train_temp = []
    cv_temp = []

```

```

for j in max_depth:
    rf = RandomForestClassifier(n_estimators=i,max_depth=j,class_weight="balanced")
    rf.fit(set1,y_train)

    y_train_pred = rf.predict(set1)
    y_cv_pred = rf.predict(set1_v)


    # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the
    positive class
    # not the predicted outputs
    train_temp.append(roc_auc_score(y_train,y_train_pred))
    cv_temp.append(roc_auc_score(y_val, y_cv_pred))
print("Completed = {}".format(i))
train_auc.append(train_temp)
cv_auc.append(cv_temp)

```

Completed = 5
 Completed = 10
 Completed = 50
 Completed = 100
 Completed = 200
 Completed = 500
 Completed = 1000


In [0]:

```
Plot_heatmap(train_auc,"Train Heatmap")
```



In [0]:

```
Plot_heatmap(cv_auc,"CV Heatmap")
```



Best parameter found

- n_estimators = 500
- max_depth = 2

AUC Plot

In [0]:

```
# probabilities calculation
rf = RandomForestClassifier(n_estimators=500,max_depth=2,class_weight="balanced")
rf.fit(set1,y_train)
y1_predict_prob = rf.predict_proba(set1_t)[: ,1]
y1_predict_prob_train = rf.predict_proba(set1)[: ,1]

# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
#fpr,tpr
fpr,tpr,thre = roc_curve(y_test,y1_predict_prob)


# am i doing it right here.....?
fpr_train,tpr_train,thre_train = roc_curve(y_train,y1_predict_prob_train)
```

In [0]:

```
# auc calculation for test data
roc_auc1 = metrics.auc(fpr,tpr)

# auc calculation for train data
roc_auc_train1 = metrics.auc(fpr_train,tpr_train)

# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
plt.plot(fpr,tpr,"b--",label = 'AUC test = %0.2f'%roc_auc1)
plt.plot(fpr_train,tpr_train,"y--",label = 'AUC train = %0.2f'%roc_auc_train1)
plt.title("AUC plot")
plt.xlabel("False positive rate")
plt.ylabel("True positive rate")
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0,1])
plt.ylim([0,1])
plt.legend(loc = "lower right")
plt.show()
```




Confusion Matrix

1. For Train data

In [0]:

```
y_train_pred = rf.predict(set1)
cm1 = confusion_matrix(y_train,y_train_pred)
cf_matrix(cm1,"train data")
```






1. For test data

In [0]:

```
y_test_pred = rf.predict(set1_t)
cm1 = confusion_matrix(y_test,y_test_pred)
# https://seaborn.pydata.org/generated/seaborn.heatmap.html
cf_matrix(cm1,"test data")
```



Set2 (TFIDF)

In [0]:

```
train_auc = []
cv_auc = []
heatmap = []
n_estimators = [5, 10, 50, 100, 200, 500, 1000]
max_depth = [2, 3, 4, 5, 6, 7, 8, 9, 10]

for i in n_estimators:
    train_temp = []
    cv_temp = []
    for j in max_depth:
        rf = RandomForestClassifier(n_estimators=i,max_depth=j,class_weight="balanced")
        rf.fit(set2,y_train)

        y_train_pred = rf.predict(set2)
        y_cv_pred = rf.predict(set2_v)

        # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the
        # positive class
        # not the predicted outputs
        train_temp.append(roc_auc_score(y_train,y_train_pred))
        cv_temp.append(roc_auc_score(y_val, y_cv_pred))
    print("Completed = {}".format(i))
    train_auc.append(train_temp)
    cv_auc.append(cv_temp)
```


```
Completed = 5
Completed = 10
Completed = 50
Completed = 100
Completed = 200
```

```
Completed = 500  
Completed = 1000
```

Heatmap


In [0]:

```
Plot_heatmap(train_auc, "Train Heatmap")
```



In [0]:

```
Plot_heatmap(cv_auc, "CV Heatmap")
```



Best parameter found

- n_estimators = 500
- max_depth = 2

AUC Plot

In [0]:

```
# probabilities calculation  
rf = RandomForestClassifier(n_estimators=500,max_depth=2,class_weight="balanced")  
rf.fit(set2,y_train)  
  
# probabilities calculation  
y2_predict_prob = rf.predict_proba(set2_t)[: ,1]  
y2_predict_prob_train = rf.predict_proba(set2)[: ,1]  
  
# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python  
#fpr,tpr  
fpr,tpr,thre = roc_curve(y_test,y2_predict_prob)  
  
# am i doing it right here.....?
```

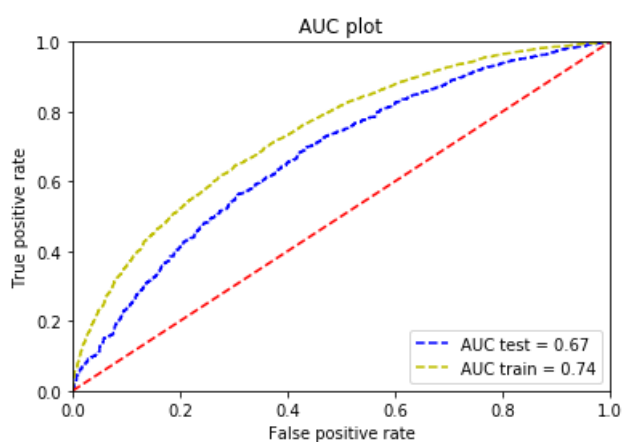
```
fpr_train,tpr_train,thre_train = roc_curve(y_train,y2_predict_prob_train)
```

In [0]:

```
# auc calculation for test data
roc_auc2 = metrics.auc(fpr,tpr)

# auc calculation for train data
roc_auc_train2 = metrics.auc(fpr_train,tpr_train)

# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
plt.plot(fpr,tpr,"b--",label = 'AUC test = %0.2f'%roc_auc2)
plt.plot(fpr_train,tpr_train,"y--",label = 'AUC train = %0.2f'%roc_auc_train2)
plt.title("AUC plot")
plt.xlabel("False positive rate")
plt.ylabel("True positive rate")
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0,1])
plt.ylim([0,1])
plt.legend(loc = "lower right")
plt.show()
```




Confusion matrix

1. Train Data

In [0]:


```
y2_train_predict = rf.predict(set2)
cm2 = confusion_matrix(y_train,y2_train_predict)
# https://seaborn.pydata.org/generated/seaborn.heatmap.html
cf_matrix(cm2,"train data")
```



1. Test data

In [0]:

```
y2_predict = rf.predict(set2_t)
cm2 = confusion_matrix(y_test,y2_predict)
# https://seaborn.pydata.org/generated/seaborn.heatmap.html
cf_matrix(cm2,"test data")
```



Set3 (Avg W2V)

In [0]:

```
train_auc = []
cv_auc = []
heatmap = []

n_estimators = [5, 10, 50, 100, 200, 500, 1000]
max_depth = [2, 3, 4, 5, 6, 7, 8, 9, 10]

for i in n_estimators:
    train_temp = []
    cv_temp = []
    for j in max_depth:
        rf = RandomForestClassifier(n_estimators=i,max_depth=j,class_weight="balanced")
        rf.fit(set3,y_train)

        y_train_pred = rf.predict(set3)
        y_cv_pred = rf.predict(set3_v)


        # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the
        # positive class
        # not the predicted outputs
        train_temp.append(roc_auc_score(y_train,y_train_pred))
        cv_temp.append(roc_auc_score(y_val, y_cv_pred))
    print("Completed = {}".format(i))
    train_auc.append(train_temp)
    cv_auc.append(cv_temp)
```


Completed = 5
Completed = 10
Completed = 50
Completed = 100
Completed = 200
Completed = 500
Completed = 1000

Heatmap

In [0]:


```
Plot_heatmap(train_auc,"Train Heatmap")
```





In [0]:

```
Plot_heatmap(cv_auc, "CV Heatmap")
```



Best Parameter found

- n_estimator = 500
- max_depth = 2

AUC Plot

In [0]:

```
# probabilities calculation
rf = RandomForestClassifier(n_estimators=500,max_depth=2,class_weight="balanced")
rf.fit(set3,y_train)

# probabilities calculation
y2_predict_prob = rf.predict_proba(set3_t)[: ,1]
y2_predict_prob_train = rf.predict_proba(set3)[: ,1]

# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
#fpr,tpr
fpr,tpr,thre = roc_curve(y_test,y2_predict_prob)

# am i doing it right here.....?
fpr_train,tpr_train,thre_train = roc_curve(y_train,y2_predict_prob_train)
```

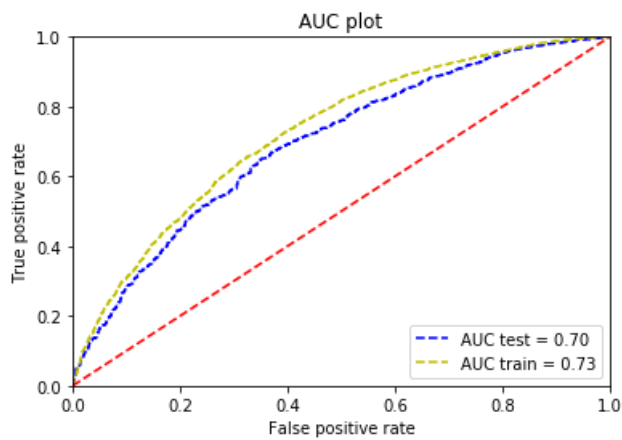
In [0]:

```
# auc calculation for test data
roc_auc3 = metrics.auc(fpr,tpr)

# auc calculation for train data
roc_auc_train3 = metrics.auc(fpr_train,tpr_train)

# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
plt.plot(fpr,tpr,"b--",label = 'AUC test = %0.2f'%roc_auc3)
plt.plot(fpr_train,tpr_train,"y--",label = 'AUC train = %0.2f'%roc_auc_train3)
```

```
plt.title("AUC plot")
plt.xlabel("False positive rate")
plt.ylabel("True positive rate")
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0,1])
plt.ylim([0,1])
plt.legend(loc = "lower right")
plt.show()
```




confusion matrix

1. Train data

In [0]:


```
y3_train_predict = rf.predict(set3)
cm3 = confusion_matrix(y_train,y3_train_predict)
# https://seaborn.pydata.org/generated/seaborn.heatmap.html
cf_matrix(cm3,"train data")
```




1. Test data

In [0]:

```
y3_predict = rf.predict(set3_t)
cm3 = confusion_matrix(y_test,y3_predict)
# https://seaborn.pydata.org/generated/seaborn.heatmap.html
cf_matrix(cm3,"test data")
```





Set4 (TFIDF Avg_w2v)

In [0]:

```
train_auc = []
cv_auc = []
heatmap = []
n_estimators = [5, 10, 50, 100, 200, 500, 1000]
max_depth = [2, 3, 4, 5, 6, 7, 8, 9, 10]

for i in n_estimators:
    train_temp = []
    cv_temp = []
    for j in max_depth:
        rf = RandomForestClassifier(n_estimators=i,max_depth=j,class_weight="balanced")
        rf.fit(set4,y_train)

        y_train_pred = rf.predict(set4)
        y_cv_pred = rf.predict(set4_v)


        # roc_auc_score(y_true, y_score) the 2nd parameter should be probability estimates of the
        # positive class
        # not the predicted outputs
        train_temp.append(roc_auc_score(y_train,y_train_pred))
        cv_temp.append(roc_auc_score(y_val, y_cv_pred))
    print("Completed = {}".format(i))
    train_auc.append(train_temp)
    cv_auc.append(cv_temp)
```

Completed = 5
 Completed = 10
 Completed = 50
 Completed = 100
 Completed = 200
 Completed = 500
 Completed = 1000

Heatmap


In [0]:

```
Plot_heatmap(train_auc,"Train Heatmap")
```



In [0]:

```
Plot_heatmap(cv_auc,"CV Heatmap")
```



AUC Plot

In [0]:

```
# probabilities calculation
rf = RandomForestClassifier(n_estimators=500,max_depth=2,class_weight="balanced")
rf.fit(set4,y_train)

# probabilities calculation
y2_predict_prob = rf.predict_proba(set4_t)[: ,1]
y2_predict_prob_train = rf.predict_proba(set4)[: ,1]

# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
#fpr,tpr
fpr,tpr,thre = roc_curve(y_test,y2_predict_prob)


# am i doing it right here.....?
fpr_train,tpr_train,thre_train = roc_curve(y_train,y2_predict_prob_train)
```


In [0]:

```
# auc calculation for test data
roc_auc4 = metrics.auc(fpr,tpr)

# auc calculation for train data
roc_auc_train4 = metrics.auc(fpr_train,tpr_train)

# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
plt.plot(fpr,tpr,"b--",label = 'AUC test = %0.2f'%roc_auc4)
plt.plot(fpr_train,tpr_train,"y--",label = 'AUC train = %0.2f'%roc_auc_train4)
plt.title("AUC plot")
plt.xlabel("False positive rate")
plt.ylabel("True positive rate")
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0,1])
plt.ylim([0,1])
plt.legend(loc = "lower right")
plt.show()
```






Confusion Matrix

1. Train Data

In [0]:


```
y4_train_predict = rf.predict(set4)
cm1 = confusion_matrix(y_train,y4_train_predict)
# https://seaborn.pydata.org/generated/seaborn.heatmap.html
cf_matrix(cm1,"train data")
```



1. Test Data

In [0]:

```
y4_predict = rf.predict(set4_t)
cm1 = confusion_matrix(y_test,y4_predict)
# https://seaborn.pydata.org/generated/seaborn.heatmap.html
cf_matrix(cm1,"test data")
```



Observation

=====

1. TFIDF found to be best Vectorizer
2. Best **train AUC of 0.72** and **test auc of 0.83** was found in case of TFIDF
3. Overfitting in model found for higher values of max_depth

Summary of RF

In [0]:

```
from prettytable import PrettyTable
summary = PrettyTable()
```

In [0]:

```
summary.field_names = ["Set", "Vectorizer", "Model", "Hyperparameter", "Test", "Train"]
```

In [0]:

```
summary.add_row(["set1", "BOW", "RF", "max_depth': 2, 'n_estimators': 500", 0.68, 0.73])
summary.add_row(["set2", "TFIDF", "RF", "'max_depth': 2, 'n_estimators': 500", 0.67, 0.74])
summary.add_row(["set3", "Avg-W2v", "RF", "'max_depth': 2, 'n_estimators': 500", 0.70, 0.73])
summary.add_row(["set4", "TFIDF W2V", "RF", "'max_depth': 2, 'n_estimators': 500", 0.69, 0.72])

print(summary)
```

Set	Vectorizer	Model	Hyperparameter	Test	Train
set1	BOW	RF	max_depth': 2, 'n_estimators': 500	0.68	0.73
set2	TFIDF	RF	'max_depth': 2, 'n_estimators': 500	0.67	0.74
set3	Avg-W2v	RF	'max_depth': 2, 'n_estimators': 500	0.7	0.73
set4	TFIDF W2V	RF	'max_depth': 2, 'n_estimators': 500	0.69	0.72

GBDT

- Tuned over wide range of below parameters of XGBOOST.
 - n_estimators = [5, 10, 50, 100, 200, 500, 1000]
 - max_depth = [2, 3, 4, 5, 6, 7, 8, 9, 10]

Using all other as default

In [0]:

```
n_estimators = [5, 10, 50, 100, 200, 500, 1000]
max_depth = [2, 3, 4, 5, 6, 7, 8, 9, 10]

def Plot_heatmap(mean_score, name):
    """
    This function plots heatmap.
    """
    df = pd.DataFrame(mean_score, index = n_estimators, columns = max_depth)
    sns.heatmap(df, annot = True)
    plt.ylabel("n_estimators")
    plt.xlabel("max_depth")
    plt.title(name)
    plt.show()
```

Set1 BOW

In [0]:

```
train_auc = []
cv_auc = []

n_estimators = [5, 10, 50, 100, 200, 500, 1000]
max_depth = [2, 3, 4, 5, 6, 7, 8, 9, 10]

for i in n_estimators:
```

```

train_temp = []
cv_temp = []
for j in max_depth:
    xgb = XGBClassifier(max_depth=j,n_estimators=i)
    xgb.fit(set1,y_train)

    y_train_pred = xgb.predict(set1)
    y_cv_pred = xgb.predict(set1_v)

    train_temp.append(roc_auc_score(y_train,y_train_pred))
    cv_temp.append(roc_auc_score(y_val, y_cv_pred))
print("Completed = {}".format(i))
train_auc.append(train_temp)
cv_auc.append(cv_temp)

```

```


Completed = 5
Completed = 10
Completed = 50
Completed = 100
Completed = 200
Completed = 500
Completed = 1000

```

Heatmap


In [0]:

```
Plot_heatmap(train_auc,"Train Heatmap")
```



In [0]:

```
Plot_heatmap(cv_auc,"CV Heatmap")
```



Best paramter found:

```
max_depth = 4
```


- max_depth = 4
- n_estimators = 100

AUC Plot

In [0]:

```
# probabilities calculation
xgb = XGBClassifier(max_depth=4,min_child_weight=4,n_estimators=100)
xgb.fit(set1,y_train)
y1_predict_prob = xgb.predict_proba(set1_t)[: ,1]
y1_predict_prob_train = xgb.predict_proba(set1)[: ,1]

# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
#fpr,tpr
fpr,tpr,thre = roc_curve(y_test,y1_predict_prob)

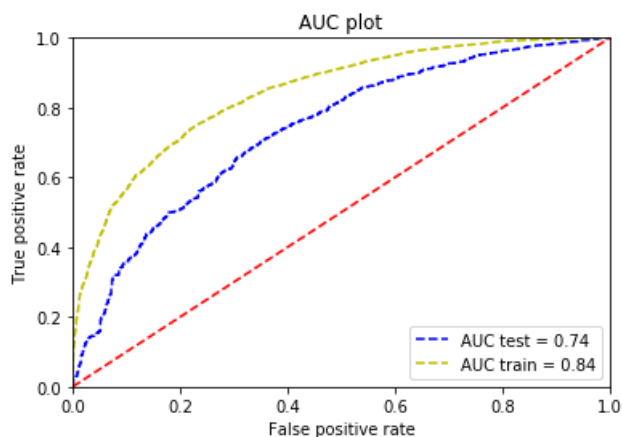
# am i doing it right here.....?
fpr_train,tpr_train,thre_train = roc_curve(y_train,y1_predict_prob_train)
```

In [0]:

```
# auc calculation for test data
roc_auc1 = metrics.auc(fpr,tpr)

# auc calculation for train data
roc_auc_train1 = metrics.auc(fpr_train,tpr_train)

# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
plt.plot(fpr,tpr,"b--",label = 'AUC test = %0.2f'%roc_auc1)
plt.plot(fpr_train,tpr_train,"y--",label = 'AUC train = %0.2f'%roc_auc_train1)
plt.title("AUC plot")
plt.xlabel("False positive rate")
plt.ylabel("True positive rate")
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0,1])
plt.ylim([0,1])
plt.legend(loc = "lower right")
plt.show()
```




Confusion Matrix


1. Train Data

In [0]:

```
y_train_pred = xgb.predict(set1)
cm1 = confusion_matrix(y_train,y_train_pred)
cf_matrix(cm1,"train data")
```

Confusion Matrix for train data






1. Text Data

In [0]:

```
y_test_pred = xgb.predict(set1_t)
cm1 = confusion_matrix(y_test, y_test_pred)
# https://seaborn.pydata.org/generated/seaborn.heatmap.html
cf_matrix(cm1, "test data")
```



Set2 (TFIDF)

In [0]:

```
train_auc = []
cv_auc = []

n_estimators = [5, 10, 50, 100, 200, 500, 1000]
max_depth = [2, 3, 4, 5, 6, 7, 8, 9, 10]

for i in n_estimators:
    train_temp = []
    cv_temp = []
    for j in max_depth:
        xgb = XGBClassifier(max_depth=j, n_estimators=i)
        xgb.fit(set2, y_train)

        y_train_pred = xgb.predict(set2)
        y_cv_pred = xgb.predict(set2_v)

        train_temp.append(roc_auc_score(y_train, y_train_pred))
        cv_temp.append(roc_auc_score(y_val, y_cv_pred))
    print("Completed = {}".format(i))
    train_auc.append(train_temp)
    cv_auc.append(cv_temp)
```


Completed = 5
 Completed = 10
 Completed = 50
 Completed = 100

Completed = 200
Completed = 500
Completed = 1000

Heatmap


In [0]:

```
Plot_heatmap(train_auc, "Train Heatmap")
```



In [0]:

```
Plot_heatmap(cv_auc, "CV Heatmap")
```



Best paramter found:

- max_depth = 4
- n_estimators = 100

AUC Plot

In [0]:

```
# probabilities calculation
xgb = XGBClassifier(max_depth=4,min_child_weight=4,n_estimators=100)
xgb.fit(set2,y_train)
y1_predict_prob = xgb.predict_proba(set2_t)[: ,1]
y1_predict_prob_train = xgb.predict_proba(set2)[: ,1]

# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
#fpr,tpr
fpr,tpr,thre = roc_curve(y_test,y1_predict_prob)


# am i doing it right here.....?
fpr_train,tpr_train,thre_train = roc_curve(y_train,y1_predict_prob_train)
```

In [0]:

```
# auc calculation for test data
roc_auc2 = metrics.auc(fpr, tpr)

# auc calculation for train data
roc_auc_train2 = metrics.auc(fpr_train, tpr_train)

# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
plt.plot(fpr, tpr, "b--", label = 'AUC test = %0.2f'%roc_auc2)
plt.plot(fpr_train, tpr_train, "y--", label = 'AUC train = %0.2f'%roc_auc_train2)
plt.title("AUC plot")
plt.xlabel("False positive rate")
plt.ylabel("True positive rate")
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.legend(loc = "lower right")
plt.show()
```




Confusion Matrix

1. Train Data

In [0]:

```
y2_train_predict = xgb.predict(set2)
cm2 = confusion_matrix(y_train, y2_train_predict)
# https://seaborn.pydata.org/generated/seaborn.heatmap.html
cf_matrix(cm2, "train data")
```




1. Text Data

In [0]:

```

y2_predict = xgb.predict(set2_t)
cm2 = confusion_matrix(y_test,y2_predict)
# https://seaborn.pydata.org/generated/seaborn.heatmap.html
cf_matrix(cm2,"test data")

```



SET3

- Reducing `n_estimators` and `max_depth` as using all leading to crash of programme

In [45]:

```

train_auc = []
cv_auc = []

n_estimators = [50, 100, 200, 500]
max_depth = [2, 3, 4, 5, 6, 7, 8]

for i in n_estimators:
    train_temp = []
    cv_temp = []
    for j in max_depth:
        xgb = XGBClassifier(max_depth=j, n_estimators=i)
        xgb.fit(set3, y_train)

        y_train_pred = xgb.predict(set3)
        y_cv_pred = xgb.predict(set3_v)

        train_temp.append(roc_auc_score(y_train, y_train_pred))
        cv_temp.append(roc_auc_score(y_val, y_cv_pred))
    print("Completed = {}".format(i))
    train_auc.append(train_temp)
    cv_auc.append(cv_temp)

```

```

Completed = 50
Completed = 100
Completed = 200
Completed = 500


```

Heatmap

In [46]:


```
Plot_heatmap(train_auc,"Train Heatmap")
```





In [47]:

```
Plot_heatmap(cv_auc, "CV Heatmap")
```



Best paramter found:

- max_depth = 3
- min_child_weight = 4
- n_estimators = 100

AUC Plot

In [0]:

```
# probabilities calculation
xgb = XGBClassifier(max_depth=3,min_child_weight=4,n_estimators=100)
xgb.fit(set3,y_train)
y1_predict_prob = xgb.predict_proba(set3_t)[: ,1]
y1_predict_prob_train = xgb.predict_proba(set3)[: ,1]

# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
#fpr,tpr
fpr,tpr,thre = roc_curve(y_test,y1_predict_prob)

# am i doing it right here.....?
fpr_train,tpr_train,thre_train = roc_curve(y_train,y1_predict_prob_train)
```


In [52]:

```
# auc calculation for test data
roc_auc3 = metrics.auc(fpr,tpr)

# auc calculation for train data
roc_auc_train3 = metrics.auc(fpr_train,tpr_train)

# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
plt.plot(fpr,tpr,"b--",label = 'AUC test = %0.2f'%roc_auc3)
plt.plot(fpr_train,tpr_train,"y--",label = 'AUC train = %0.2f'%roc_auc_train3)
plt.title("AUC plot")
plt.xlabel("False positive rate")
plt.ylabel("True positive rate")
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0,1])
```

```
plt.ylim([0,1])
plt.legend(loc = "lower right")
plt.show()
```




Confusion Matrix

1. Train Data

In [53]:


```
y3_train_predict = xgb.predict(set3)
cm3 = confusion_matrix(y_train,y3_train_predict)
# https://seaborn.pydata.org/generated/seaborn.heatmap.html
cf_matrix(cm3,"train data")
```



1. Test Data

In [54]:

```
y3_predict = xgb.predict(set3_t)
cm3 = confusion_matrix(y_test,y3_predict)
# https://seaborn.pydata.org/generated/seaborn.heatmap.html
cf_matrix(cm3,"test data")
```





Set 4

- Reducing `n_estimators` and `max_depth` as using all leading to crash of programme

In [48]:

```
train_auc = []
cv_auc = []

n_estimators = [50, 100, 200, 500]
max_depth = [2, 3, 4, 5, 6, 7, 8]

for i in n_estimators:
    train_temp = []
    cv_temp = []
    for j in max_depth:
        xgb = XGBClassifier(max_depth=j, n_estimators=i)
        xgb.fit(set4, y_train)

        y_train_pred = xgb.predict(set4)
        y_cv_pred = xgb.predict(set4_v)


        train_temp.append(roc_auc_score(y_train, y_train_pred))
        cv_temp.append(roc_auc_score(y_val, y_cv_pred))
    print("Completed = {}".format(i))
    train_auc.append(train_temp)
    cv_auc.append(cv_temp)
```

Completed = 50
 Completed = 100
 Completed = 200
 Completed = 500

Heatmap


In [49]:


```
Plot_heatmap(train_auc, "Train Heatmap")
```



In [50]:

```
Plot_heatmap(cv_auc, "CV Heatmap")
```





Best paramter found:

- max_depth = 3
- n_estimators = 100

AUC Plot

In [0]:

```
# probabilities calculation
xgb = XGBClassifier(max_depth=3,in_child_weight=4,n_estimators=100)
xgb.fit(set4,y_train)
y1_predict_prob = xgb.predict_proba(set4_t)[: ,1]
y1_predict_prob_train = xgb.predict_proba(set4)[: ,1]

# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
#fpr,tpr
fpr,tpr,thre = roc_curve(y_test,y1_predict_prob)


# am i doing it right here.....?
fpr_train,tpr_train,thre_train = roc_curve(y_train,y1_predict_prob_train)
```

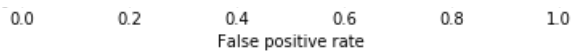
In [56]:

```
# auc calculation for test data
roc_auc4 = metrics.auc(fpr,tpr)

# auc calculation for train data
roc_auc_train4 = metrics.auc(fpr_train,tpr_train)

# took reference from https://stackoverflow.com/questions/25009284/how-to-plot-roc-curve-in-python
plt.plot(fpr,tpr,"b--",label = 'AUC test = %0.2f'%roc_auc4)
plt.plot(fpr_train,tpr_train,"y--",label = 'AUC train = %0.2f'%roc_auc_train4)
plt.title("AUC plot")
plt.xlabel("False positive rate")
plt.ylabel("True positive rate")
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0,1])
plt.ylim([0,1])
plt.legend(loc = "lower right")
plt.show()
```






Confusion Matrix

1. Train Data

In [58]:


```
y4_train_predict = xgb.predict(set4)
cm1 = confusion_matrix(y_train,y4_train_predict)
# https://seaborn.pydata.org/generated/seaborn.heatmap.html
cf_matrix(cm1,"train data")
```



1. Test Data

In [59]:

```
y4_predict = xgb.predict(set4_t)
cm1 = confusion_matrix(y_test,y4_predict)
# https://seaborn.pydata.org/generated/seaborn.heatmap.html
cf_matrix(cm1,"test data")
```



Obseravtion:

=====

1. **max_depth** found to be most important parameter in tuning xgboost
2. For given max_depth and min_child_weight increase in n_estimators was leading to overfitting the model.
3. Tuning Avg-W2V and TFIDF-W2V took more time as compared to others.

Summary

Summary

In [0]:

```
from prettytable import PrettyTable
summary = PrettyTable()
```

In [0]:

```
summary.field_names = ["Set", "Vectorizer", "Model", "Hyperparameter", "Test", "Train"]
```

In [62]:

```
summary.add_row(["set1", "BOW", "XGBOOST", "max_depth = 4,min_child_weight = 4,n_estimators = 100", 0.74, 0.84])
summary.add_row(["set2", "TFIDF", "XGBOOST", "max_depth=4,min_child_weight = 4,n_estimators=100", 0.74, 0.86])
summary.add_row(["set3", "Avg-W2v", "XGBOOST", "max_depth=3,min_child_weight = 4,n_estimators=100", 0.72, 0.84])
summary.add_row(["set4", "TFIDF W2V", "XGBOOST", "max_depth=3,min_child_weight = 4,n_estimators=100", 0.72, 0.84])

print(summary)
```

Set	Vectorizer	Model	Hyperparameter	Test	Train
set1	BOW	XGBOOST	max_depth = 4,min_child_weight = 4,n_estimators = 100	0.74	0.84
set2	TFIDF	XGBOOST	max_depth=4,min_child_weight = 4,n_estimators=100	0.74	0.86
set3	Avg-W2v	XGBOOST	max_depth=3,min_child_weight = 4,n_estimators=100	0.72	0.84
set4	TFIDF W2V	XGBOOST	max_depth=3,min_child_weight = 4,n_estimators=100	0.72	0.84

In [0]:

In [0]: